

Terry Stop Analysis

In the 1968 Supreme Court case "Terry v. Ohio", the court found that a police officer was not in violation of the "unreasonable search and seizure" clause of the Fourth Amendment after he stopped and frisked suspects only because their behavior was suspicious. Thus the phrase "Terry Stops" are in reference to stops made of suspicious drivers.

This is an analysis of over 48,000 Terry Stops, with a goal of predicting if an arrest will be made based off time of day, whether a suspect was frisked, and racial & gender demographics of both the suspects and officers.

The overall goal of the analysis is to have the highest possible recall, to minimize false positives, accidentally classifying subjects who were not arrested as arrested.

In [1]: *import the necessary libraries.*

```
import pandas as pd
import numpy as np
import seaborn as sns
import datetime
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import plot_confusion_matrix
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import precision_score, recall_score, accuracy_score
from xgboost import XGBClassifier
from sklearn.tree import DecisionTreeClassifier
```

Step 1: Cleaning the Data

```
In [2]: # Import and look at the rows of our dataset.

pd.set_option('display.max_columns', None)
df = pd.read_csv('Terry_Stops.csv')
df.head()
```

Out[2]:

	Subject Age Group	Subject ID	GO / SC Num	Terry Stop ID	Stop Resolution	Weapon Type	Officer ID	Officer YOB	Officer Gender	Of I
0	-	-1	20140000120677	92317	Arrest	None	7500	1984	M	Bla Af Ame
1	-	-1	20150000001463	28806	Field Contact	None	5670	1965	M	v
2	-	-1	20150000001516	29599	Field Contact	None	4844	1961	M	v
3	-	-1	20150000001670	32260	Field Contact	None	7539	1963	M	v
4	-	-1	20150000001739	33155	Field Contact	None	6973	1977	M	v

```
In [3]: # For ease of analysis, let's add underscores to each column name
renamed_columns = ['Subject_Age_Group', 'Subject_ID', 'GO_SC_Num', 'Terry
                  'Stop_Resolution', 'Weapon_Type', 'Officer_ID', 'Offi
                  'Officer_Gender', 'Officer_Race', 'Subject_Perceived_R
                  'Subject_Perceived_Gender', 'Reported_Date', 'Reported
                  'Initial_Call_Type', 'Final_Call_Type', 'Call_Type',
                  'Arrest_Flag', 'Frisk_Flag', 'Precinct', 'Sector', 'Bea
df.columns = renamed_columns
df.head()
```

Out[3]:

	Subject_Age_Group	Subject_ID	GO_SC_Num	Terry_Stop_ID	Stop_Resolution	Weapon_Typ
0	-	-1	20140000120677	92317	Arrest	Non
1	-	-1	20150000001463	28806	Field Contact	Non
2	-	-1	20150000001516	29599	Field Contact	Non
3	-	-1	20150000001670	32260	Field Contact	Non
4	-	-1	20150000001739	33155	Field Contact	Non

```
In [4]: # Looking at the initial data, 'Subject_ID, GO_SC_Num', 'Terry_Stop_ID',
# individual indentifiers so it's okay to drop those
```

```
In [5]: df = df.drop(['Subject_ID', 'GO_SC_Num', 'Terry_Stop_ID', 'Officer_ID'],
```

In [6]: *# Take a look to see if there are any null values in the data*

```
df.isna().sum()
```

```
Out[6]: Subject_Age_Group      0
        Stop_Resolution      0
        Weapon_Type          0
        Officer_YOB          0
        Officer_Gender       0
        Officer_Race         0
        Subject_Perceived_Race 0
        Subject_Perceived_Gender 0
        Reported_Date        0
        Reported_Time        0
        Initial_Call_Type     0
        Final_Call_Type      0
        Call_Type            0
        Officer_Squad        605
        Arrest_Flag          0
        Frisk_Flag           0
        Precinct             0
        Sector               0
        Beat                 0
        dtype: int64
```

In [7]: *# Interesting, because upon initial analysis, there are empty entries
in 'Subject Age Group', all of the 'Call Type' columns*

```
df.info()
```

```
Stop_Resolution      48094 non-null object
Weapon_Type          48094 non-null object
Officer_YOB          48094 non-null int64
Officer_Gender       48094 non-null object
Officer_Race         48094 non-null object
Subject_Perceived_Race 48094 non-null object
Subject_Perceived_Gender 48094 non-null object
Reported_Date        48094 non-null object
Reported_Time        48094 non-null object
Initial_Call_Type     48094 non-null object
Final_Call_Type      48094 non-null object
Call_Type            48094 non-null object
Officer_Squad        47489 non-null object
Arrest_Flag          48094 non-null object
Frisk_Flag           48094 non-null object
Precinct             48094 non-null object
Sector               48094 non-null object
Beat                 48094 non-null object
dtypes: int64(1), object(18)
memory usage: 7.0+ MB
```

In [8]: *# investigate the columns that are mostly '-' to see if there is data
or if it's mostly null values*

```
In [9]: print('Subject Age Group:' , '\n', df['Subject_Age_Group'].value_counts()
print('Initial Call Type:' , '\n', df['Initial_Call_Type'].value_counts())
print('Final Call Type:' , '\n', df['Final_Call_Type'].value_counts(), '
print('Call Type:' , "\n", df['Call_Type'].value_counts(), '\n')
```

Subject Age Group:

26 - 35	16012
36 - 45	10262
18 - 25	9590
46 - 55	6225
56 and Above	2448
1 - 17	1986
-	1571

Name: Subject_Age_Group, dtype: int64

Initial Call Type:

-	13207
SUSPICIOUS STOP - OFFICER INITIATED ONVIEW	3221
SUSPICIOUS PERSON, VEHICLE OR INCIDENT	3078
DISTURBANCE, MISCELLANEOUS/OTHER	2477
ASLT - IP/JO - WITH OR W/O WPNS (NO SHOOTINGS)	2070

...

PREPLANNED EVENT	1
-ASSIGNED DUTY - STAKEOUT	1
ALARM - RESIDENTIAL - SILENT/AUD PANIC/DURESS	1
WARRANT PICKUP - FROM OTHER AGENCY	1
ORDER - ASSIST DV VIC W/SRVC OF COURT ORDER	1

Name: Initial_Call_Type, Length: 168, dtype: int64

Final Call Type:

-	13207
--SUSPICIOUS CIRCUM. - SUSPICIOUS PERSON	3873
--PROWLER - TRESPASS	3381
--DISTURBANCE - OTHER	2775
--ASSAULTS, OTHER	2344

...

SHOTS -DELAY/INCLUDES HEARD/NO ASSAULT	1
MVC - UNK INJURIES	1
MVC - REPORT, NON INJ/NON BLKG OR AFTER FACT INJ	1
FIGHT - VERBAL/ORAL (NO WEAPONS)	1
ESCAPE - PRISONER	1

Name: Final_Call_Type, Length: 211, dtype: int64

Call Type:

911	21844
-	13207
ONVIEW	9280
TELEPHONE OTHER, NOT 911	3399
ALARM CALL (NOT POLICE ALARM)	355
TEXT MESSAGE	8
SCHEDULED EVENT (RECURRING)	1

Name: Call_Type, dtype: int64

```
In [10]: 1 "Call Type" is mostly unnecessary and can be dropped. Also 'Initial_Cal  
missing over 13k columns so it's okay to drop those
```

```
In [11]: df = df.drop(['Call_Type', 'Initial_Call_Type', 'Final_Call_Type'], axis=
```

```
In [12]: df['Subject_Age_Group'].value_counts()
```

```
Out[12]: 26 - 35          16012  
         36 - 45          10262  
         18 - 25           9590  
         46 - 55           6225  
         56 and Above      2448  
         1 - 17           1986  
         -                1571  
         Name: Subject_Age_Group, dtype: int64
```

```
In [13]: df = df[df.Subject_Age_Group != '-']
```

```
In [14]: df['Subject_Age_Group'].value_counts()
```

```
Out[14]: 26 - 35          16012  
         36 - 45          10262  
         18 - 25           9590  
         46 - 55           6225  
         56 and Above      2448  
         1 - 17           1986  
         Name: Subject_Age_Group, dtype: int64
```

```
In [15]: df['Officer_YOB'].value_counts()
```

```
Out[15]: 1986      3261
          1987      2934
          1984      2691
          1991      2659
          1985      2440
          1992      2426
          1990      2274
          1988      2114
          1989      2010
          1982      1833
          1983      1681
          1979      1493
          1993      1482
          1981      1406
          1995      1232
          1971      1182
          1978      1137
          1977       994
          1976       993
          1994       905
          1973       904
          1980       809
          1967       701
          1996       687
          1968       590
          1970       559
          1969       539
          1974       533
          1975       521
          1997       451
          1962       449
          1964       431
          1972       413
          1965       412
          1963       236
          1958       215
          1961       206
          1966       178
          1959       167
          1960       128
          1954        44
          1957        43
          1998        36
          1953        33
          1900        31
          1955        21
          1956        17
          1948        10
          1949         5
          1952         4
          1946         2
          1951         1
          Name: Officer_YOB, dtype: int64
```

```
In [16]: def officer_yob_decade(x):
         if (x <= 1959):
             return '1900-1960'
         elif (x > 1959) and (x <= 1969):
             return '1960s'
         elif (x > 1969) and (x <= 1979):
             return '1970s'
         elif (x > 1979) and (x <= 1989):
             return '1980s'
         elif (x > 1989) and (x <= 1999):
             return '1990s'
```

```
In [17]: df['Officer_Age_By_Decade'] = df['Officer_YOB'].apply(officer_yob_decade)
         df['Officer_Age_By_Decade'].value_counts()
```

```
Out[17]: 1980s      21179
         1990s      12152
         1970s       8729
         1960s       3870
         1900-1960    593
         Name: Officer_Age_By_Decade, dtype: int64
```

```
In [18]: # Explore "Arrest Flag" and "Stop Resolution"
         # as they both have arrest data that will serve
         # as the target variable for this exploration.
```

```
In [19]: df['Arrest_Flag'].value_counts()
```

```
Out[19]: N      42843
         Y       3680
         Name: Arrest_Flag, dtype: int64
```

```
In [20]: df['Stop_Resolution'].value_counts()
```

```
Out[20]: Field Contact      18637
         Offense Report    15287
         Arrest            11706
         Referred for Prosecution    719
         Citation / Infraction    174
         Name: Stop_Resolution, dtype: int64
```

```
In [21]: # It seems there is a discrepancy in Arrest data between 'Stop Resolution'
         # and "Arrest_Flag". I am making the executive decision to base this
         # exploration around "Stop Resolution" as it reads as more thorough
         # in it's reporting of resolution rather than a simple 'Yes'/'No' in
         # 'Arrest Flag'.
```

```
In [22]: df['Stop_Resolution'] = df['Stop_Resolution'].apply(lambda x: 'Yes' if x
         df['Stop_Resolution'].value_counts())
```

```
Out[22]: No      34817
         Yes      11706
         Name: Stop_Resolution, dtype: int64
```



```
In [23]: df = df.drop('Arrest_Flag', axis=1)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 46523 entries, 124 to 48093
Data columns (total 16 columns):
Subject_Age_Group      46523 non-null object
Stop_Resolution        46523 non-null object
Weapon_Type           46523 non-null object
Officer_YOB            46523 non-null int64
Officer_Gender         46523 non-null object
Officer_Race           46523 non-null object
Subject_Perceived_Race 46523 non-null object
Subject_Perceived_Gender 46523 non-null object
Reported_Date          46523 non-null object
Reported_Time          46523 non-null object
Officer_Squad          45936 non-null object
Frisk_Flag             46523 non-null object
Precinct               46523 non-null object
Sector                 46523 non-null object
Beat                   46523 non-null object
Officer_Age_By_Decade  46523 non-null object
dtypes: int64(1), object(15)
memory usage: 6.0+ MB
```

```
In [24]: # Now explore factors like the races and genders of both the officers
         # and the people pulled over
```

```
In [25]: ('Officer Gender:' , '\n', df['Officer_Gender'].value_counts(), '\n')
('Officer Race:', '\n', df['Officer_Race'].value_counts(), '\n')
('Subject Perceived Gender:', '\n', df['Subject_Perceived_Gender'].valu
('Subject Perceived Race:', "\n", df['Subject_Perceived_Race'].value_cou
```

Officer Gender:

M	41154
F	5344
N	25

Name: Officer_Gender, dtype: int64

Officer Race:

White	35173
Hispanic or Latino	2679
Two or More Races	2669
Asian	1987
Black or African American	1790
Not Specified	1447
Nat Hawaiian/Oth Pac Islander	437
American Indian/Alaska Native	310
Unknown	31

Name: Officer_Race, dtype: int64

Subject Perceived Gender:

Male	36887
Female	9489
Unable to Determine	108
Unknown	17
-	16
Gender Diverse (gender non-conforming and/or transgender)	6

Name: Subject_Perceived_Gender, dtype: int64

Subject Perceived Race:

White	23078
Black or African American	13943
Unknown	2466
Hispanic	1659
Asian	1540
-	1471
American Indian or Alaska Native	1352
Multi-Racial	801
Other	150
Native Hawaiian or Other Pacific Islander	63

Name: Subject_Perceived_Race, dtype: int64

```
In [26]: # There seem to be many different ways to say "unknown" in these columns
# Let's combine the redundant values together
```

```
In [27]: 'Officer_Race' = df['Officer_Race'].apply(lambda x: "Other" if x in ['N
'Subject_Perceived_Gender' = df['Subject_Perceived_Gender'].apply(lambda
'Subject_Perceived_Race' = df['Subject_Perceived_Race'].apply(lambda x:
```

```
In [28]: print(df['Officer_Race'].value_counts(), '\n')
print(df['Subject_Perceived_Gender'].value_counts(), '\n')
print(df['Subject_Perceived_Race'].value_counts(), '\n')
```

White	35173
Hispanic or Latino	2679
Two or More Races	2669
Asian	1987
Black or African American	1790
Other	1478
Nat Hawaiian/Oth Pac Islander	437
American Indian/Alaska Native	310

Name: Officer_Race, dtype: int64

Male	36887
Female	9489
Unknown/GNC	147

Name: Subject_Perceived_Gender, dtype: int64

White	23078
Black or African American	13943
Unknown	4087
Hispanic	1659
Asian	1540
American Indian or Alaska Native	1352
Multi-Racial	801
Native Hawaiian or Other Pacific Islander	63

Name: Subject_Perceived_Race, dtype: int64

In [29]: *# Also let's clean up the gender columns of both the officer and the
subject. Entries outside of the gender binary are miniscule compared
to Male and Female.*

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 46523 entries, 124 to 48093
Data columns (total 16 columns):
Subject_Age_Group      46523 non-null object
Stop_Resolution        46523 non-null object
Weapon_Type            46523 non-null object
Officer_YOB            46523 non-null int64
Officer_Gender         46523 non-null object
Officer_Race           46523 non-null object
Subject_Perceived_Race 46523 non-null object
Subject_Perceived_Gender 46523 non-null object
Reported_Date          46523 non-null object
Reported_Time          46523 non-null object
Officer_Squad          45936 non-null object
Frisk_Flag             46523 non-null object
Precinct               46523 non-null object
Sector                 46523 non-null object
Beat                   46523 non-null object
Officer_Age_By_Decade   46523 non-null object
dtypes: int64(1), object(15)
memory usage: 6.0+ MB
```

In [30]: `df = df[df.Officer_Gender != 'N']`
`df = df[df.Subject_Perceived_Gender != 'Unknown/GNC']`
`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 46351 entries, 124 to 48093
Data columns (total 16 columns):
Subject_Age_Group      46351 non-null object
Stop_Resolution        46351 non-null object
Weapon_Type            46351 non-null object
Officer_YOB            46351 non-null int64
Officer_Gender         46351 non-null object
Officer_Race           46351 non-null object
Subject_Perceived_Race 46351 non-null object
Subject_Perceived_Gender 46351 non-null object
Reported_Date          46351 non-null object
Reported_Time          46351 non-null object
Officer_Squad          45794 non-null object
Frisk_Flag             46351 non-null object
Precinct               46351 non-null object
Sector                 46351 non-null object
Beat                   46351 non-null object
Officer_Age_By_Decade   46351 non-null object
dtypes: int64(1), object(15)
memory usage: 6.0+ MB
```

```
In [31]: df['Weapon_Type'].value_counts()
```

```
Out[31]: None                                31478
-                                             12076
Lethal Cutting Instrument                    1452
Knife/Cutting/Stabbing Instrument           671
Handgun                                     285
Firearm Other                              92
Blunt Object/Striking Implement             89
Club, Blackjack, Brass Knuckles            48
Firearm                                     43
Mace/Pepper Spray                          30
Other Firearm                              22
Firearm (unk type)                         15
Taser/Stun Gun                             10
None/Not Applicable                        9
Club                                         9
Fire/Incendiary Device                     7
Rifle                                       6
Shotgun                                    3
Automatic Handgun                          2
Personal Weapons (hands, feet, etc.)       2
Blackjack                                  1
Brass Knuckles                             1
Name: Weapon_Type, dtype: int64
```

```
In [32]: df = df.drop('Weapon_Type', axis=1)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 46351 entries, 124 to 48093
Data columns (total 15 columns):
Subject_Age_Group      46351 non-null object
Stop_Resolution        46351 non-null object
Officer_YOBB          46351 non-null int64
Officer_Gender         46351 non-null object
Officer_Race           46351 non-null object
Subject_Perceived_Race 46351 non-null object
Subject_Perceived_Gender 46351 non-null object
Reported_Date          46351 non-null object
Reported_Time          46351 non-null object
Officer_Squad          45794 non-null object
Frisk_Flag             46351 non-null object
Precinct              46351 non-null object
Sector                46351 non-null object
Beat                  46351 non-null object
Officer_Age_By_Decade  46351 non-null object
dtypes: int64(1), object(14)
memory usage: 5.7+ MB
```

```
In [33]: # Drop Weapon type, there are far more null values than there are weapon
```

```
In [34]: df['Reported_Time'].value_counts()
```

```
Out[34]: 03:09:00    50
          02:56:00    49
          19:01:00    48
          19:18:00    48
          03:13:00    48
          ..
          14:54:22     1
          00:11:30     1
          22:24:46     1
          03:07:36     1
          21:19:26     1
          Name: Reported_Time, Length: 13183, dtype: int64
```

```
In [35]: # There are 13K different times, let's create a new column that is just
         # the hours of the stops.
```

```
In [36]: df['Reported_Time'] = pd.to_datetime(df['Reported_Time'])
         df['Reported_Hour'] = df['Reported_Time'].apply(lambda x: x.hour)
```

```
In [37]: df['Reported_Hour'].value_counts()
```

```
Out[37]: 18    2830
          2     2678
          19    2545
          3     2473
          17    2472
          23    2318
          1     2318
          16    2220
          0     2111
          15    2111
          14    2040
          22    1977
          11    1908
          13    1825
          10    1808
          21    1671
          12    1557
          5     1490
          9     1442
          4     1429
          20    1411
          8     1282
          7     1247
          6     1188
          Name: Reported_Hour, dtype: int64
```

```
In [38]: def time_of_day(x):  
         if (x > 4) and (x <= 11):  
             return 'Morning'  
         elif (x > 12) and (x <= 19):  
             return 'Afternoon'  
         else:  
             return "Night"
```

```
In [39]: df['Time_of_Day'] = df['Reported_Hour'].apply(time_of_day)
```

```
In [40]: df['Time_of_Day'].value_counts()
```

```
Out[40]: Night          19943  
         Afternoon     16043  
         Morning       10365  
         Name: Time_of_Day, dtype: int64
```

```
In [41]: # Now that those are divided, let's divide them into AM and PM and  
         # create a new column. 0 is AM, 1 is PM
```

```
In [42]: df['Reported_AM_or_PM'] = df['Reported_Hour'].apply(lambda x: 0 if x < 12 else 1)  
df['Reported_AM_or_PM'] = df['Reported_AM_or_PM'].astype(int)
```

```
In [43]: df['Reported_AM_or_PM'].value_counts()
```

```
Out[43]: 1      24977  
         0      21374  
         Name: Reported_AM_or_PM, dtype: int64
```

```
In [44]: df = df.drop(['Reported_Time'], axis=1)
```

```
In [ ]:
```

```
In [45]: df['Frisk_Flag'].value_counts()
```

```
Out[45]: N      35460  
         Y      10498  
         -        393  
         Name: Frisk_Flag, dtype: int64
```

```
In [46]: df = df[df.Frisk_Flag != '-']  
df['Frisk_Flag'] = df['Frisk_Flag'].apply(lambda x: "0" if x == 'N' else "1")
```

```
In [47]: df['Frisk_Flag'].value_counts()
```

```
Out[47]: 0      35460  
         1      10498  
         Name: Frisk_Flag, dtype: int64
```

```
In [48]: # Let's explore dates
```

```
In [49]: df['Reported_Date'].value_counts()
```

```
Out[49]: 2015-10-01T00:00:00    87
         2015-09-29T00:00:00    64
         2015-05-28T00:00:00    54
         2015-08-04T00:00:00    53
         2019-04-26T00:00:00    52
         ..
         2015-03-24T00:00:00     1
         2015-04-14T00:00:00     1
         2015-05-13T00:00:00     1
         2015-03-15T00:00:00     1
         2015-05-10T00:00:00     1
         Name: Reported_Date, Length: 2296, dtype: int64
```

```
In [50]: df['Reported_Date'] = pd.DatetimeIndex(df['Reported_Date']).month
```

```
In [51]: df['Reported_Date'].value_counts()
```

```
Out[51]: 5      4818
         6      4258
         7      4025
         4      3979
         8      3943
        10      3903
         1      3763
         3      3745
         9      3571
        11      3443
         2      3261
        12      3249
         Name: Reported_Date, dtype: int64
```

```
In [52]: df['Precinct'].value_counts()
```

```
Out[52]: West      11311
         North     10158
         -         9368
         East      6077
         South     5480
         Southwest  2213
         SouthWest  1124
         Unknown   172
         OOJ       34
         FK ERROR   21
         Name: Precinct, dtype: int64
```



```
In [53]: df = df[df.Precinct != '-']  
df = df[df.Precinct != 'Unknown']  
df = df[df.Precinct != 'OOJ']  
df = df[df.Precinct != 'FK ERROR']  
df['Precinct'] = df['Precinct'].apply(lambda x: 'SouthWest' if x in ['So  
df['Precinct'].value_counts()
```

```
Out[53]: West          11311  
North          10158  
East           6077  
South          5480  
SouthWest      3337  
Name: Precinct, dtype: int64
```

```
In [54]: df['Beat'].value_counts()
```

```
Out[54]: N3          1129  
E2           1041  
K3            978  
M2            835  
M3            768  
...  
C2             90  
U3             85  
N1             76  
J2             72  
99              2  
Name: Beat, Length: 103, dtype: int64
```

```
In [55]: df['Sector'].value_counts()
```

```
Out[55]: E          2230
         M          2195
         N          2105
         K          1741
         K          1700
         B          1611
         L          1558
         D          1468
         R          1374
         F          1327
         M          1285
         S          1284
         U          1247
         D          1161
         O          1090
         J          1088
         G          1047
         C           991
         E           944
         Q           936
         W           885
         Q           821
         N           719
         F           690
         R           633
         O           617
         B           534
         S           483
         G           466
         U           454
         L           448
         W           433
         C           398
         J           394
         -             6
         Name: Sector, dtype: int64
```

```
In [56]: df['Officer_Squad'].value_counts()
```

```
Out[56]: TRAINING - FIELD TRAINING SQUAD          4055
         WEST PCT 1ST W - DAVID/MARY             1213
         WEST PCT 2ND W - D/M RELIEF              816
         SOUTHWEST PCT 2ND W - FRANK               741
         WEST PCT 1ST W - KING/QUEEN              736
         ...
         CANINE - NIGHT SQUAD                      1
         BURG/THEFT/JUV - WEST                     1
         CANINE - DAY SQUAD                         1
         COMMUNITY OUTREACH - YOUTH VIOLENCE-SCHOOLS DETAIL 1
         TRAF - MOTORCYCLE UNIT - T2 SQUAD          1
         Name: Officer_Squad, Length: 156, dtype: int64
```

```
In [57]: # Dropping 'Officer Squad', 'Beat' and 'Sector' for reasons similar to C
```

```
In [58]: df = df.drop(['Beat', 'Sector', 'Officer_Squad'], axis=1)
```

```
In [59]: df.columns
```

```
Out[59]: Index(['Subject_Age_Group', 'Stop_Resolution', 'Officer_YOB', 'Officer_Gender',  
              'Officer_Race', 'Subject_Perceived_Race', 'Subject_Perceived_Gender',  
              'Reported_Date', 'Frisk_Flag', 'Precinct', 'Officer_Age_By_Decade',  
              'Reported_Hour', 'Time_of_Day', 'Reported_AM_or_PM'],  
              dtype='object')
```

Step 2: Visualize the Data

```
In [60]: # Begin visualization of cleaned data, starting with a visualization of  
         # our target variable, Stop Resolution.
```

```
In [61]: arrests = df['Stop_Resolution']=='Yes'
non_arrests = df['Stop_Resolution']=='No'
y = df['Stop_Resolution']
num_of_arrests = df[arrests].shape[0]
num_of_nonarrests = df[non_arrests].shape[0]

print('Target Variable: Stop Resolution')
print('Total Arrests: ', num_of_arrests)
print('Total of Non-arrests: ', num_of_nonarrests)

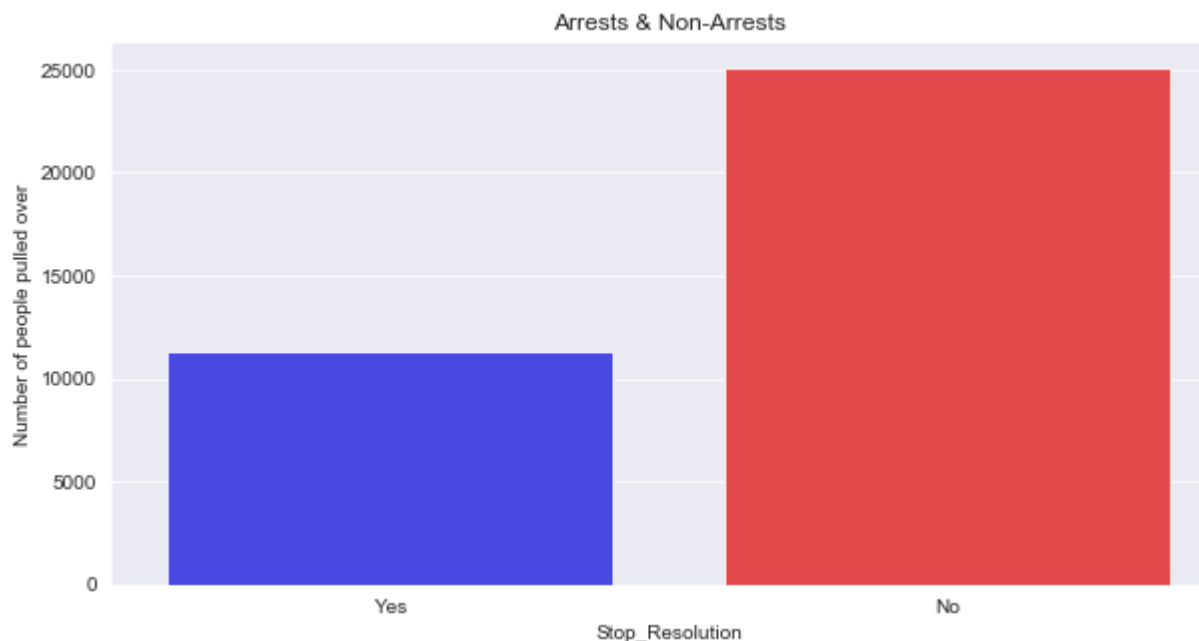
sns.set_style('darkgrid');
plt.figure(figsize = (10,5));
sns.countplot(df['Stop_Resolution'], alpha = .80, palette= ['blue', 'red'])
plt.title('Arrests & Non-Arrests');

plt.ylabel('Number of people pulled over');
plt.show()
```

Target Variable: Stop Resolution

Total Arrests: 11278

Total of Non-arrests: 25085




```
In [62]: # Create bar charts to compare different columns against the arrest dat
```

```
In [63]: subject_age_group = df.groupby(['Subject_Age_Group', 'Stop_Resolution'])
subject_age_group['% of Stops Ending In Arrest'] = (subject_age_group['Y
print('Subject Age Group\n')
print(subject_age_group)

viz_1 = subject_age_group.plot(kind = 'bar', stacked = True,
                                title = "People Pulled Over By Age: Arrested
                                color = ['blue', 'red', 'white'], alpha = .70
```

Subject Age Group

Stop_Resolution	No	Yes	% of Stops Ending In Arrest
Subject_Age_Group			
1 - 17	1204	420	25.862069
18 - 25	5155	2304	30.888859
26 - 35	8480	4132	32.762448
36 - 45	5511	2542	31.565876
46 - 55	3391	1357	28.580455
56 and Above	1344	523	28.012855




```
In [64]: subject_perceived_race = df.groupby(['Subject_Perceived_Race', 'Stop_Res
subject_perceived_race['% of Stops Ending In Arrest'] = (subject_perceiv
print('Subject Perceived Race\n')
print(subject_perceived_race)

viz_2 = subject_perceived_race.plot(kind = 'bar', stacked = True,
                                     title = 'People Pulled Over by Percei
                                     color = ['blue', 'red', 'white'], alph
```

Subject Perceived Race

Stop_Resolution	No	Yes	\
Subject_Perceived_Race			
American Indian or Alaska Native	669	386	
Asian	801	417	
Black or African American	7444	3795	
Hispanic	811	379	
Multi-Racial	404	155	
Native Hawaiian or Other Pacific Islander	41	18	
Unknown	2612	721	
White	12303	5407	

Stop_Resolution	% of Stops Ending In Arrest
Subject_Perceived_Race	
American Indian or Alaska Native	36.587678
Asian	34.236453
Black or African American	33.766349
Hispanic	31.848739
Multi-Racial	27.728086
Native Hawaiian or Other Pacific Islander	30.508475
Unknown	21.632163
White	30.530774




```
In [65]: # It appears American Indian or Alaska Native have the highest percentage
# of arrests made after a Terry stop with nearly 37%. While important to
# there were only 669 total stops of 36363 total in the data set.
# The next highest was Asian with 34.2% , with Black, white, and Native
# at around 30% each.
```

```
In [66]: er', 'Stop_Resolution']).Subject_Perceived_Gender.count().unstack()
        subject_perceived_gender['Yes'] / (subject_perceived_gender.sum(axis=1))*100

True,
by Perceived Gender:\n Arrested v Not Arrested',
te'], alpha = .80, rot = 30)
```

Subject Perceived Gender

Stop_Resolution	No	Yes	% of Stops Ending In Arrest
Subject_Perceived_Gender			
Female	5356	2151	28.653257
Male	19729	9127	31.629470




In [67]: *Male subjects were arrested 31% of the arrests while female subjects were arrested 28.65%. While the male subjects were stopped nearly 4 times the number of female, their arrest percentage was surprisingly close*


```
In [68]: der = df.groupby(['Officer_Gender', 'Stop_Resolution']).Officer_Gender.count()
der['% of Stops Ending In Arrest'] = (officer_gender['Yes'] / (officer_gender['No'] + officer_gender['Yes'])) * 100
officer_gender.plot(kind = 'bar', stacked = True,
                    title = 'Officers by Gender:\n Arrested v Not Arrested',
                    color = ['blue', 'red', 'white'], alpha = .80, rot = 45)
```

Officer Gender

Stop_Resolution	No	Yes	% of Stops Ending In Arrest
Officer_Gender			
F	2972	1177	28.368282
M	22113	10101	31.355932



```
In [69]: # Male cops arrested individuals 31% of those pulled over, while male of
# arrested at a rate of 28%.
```

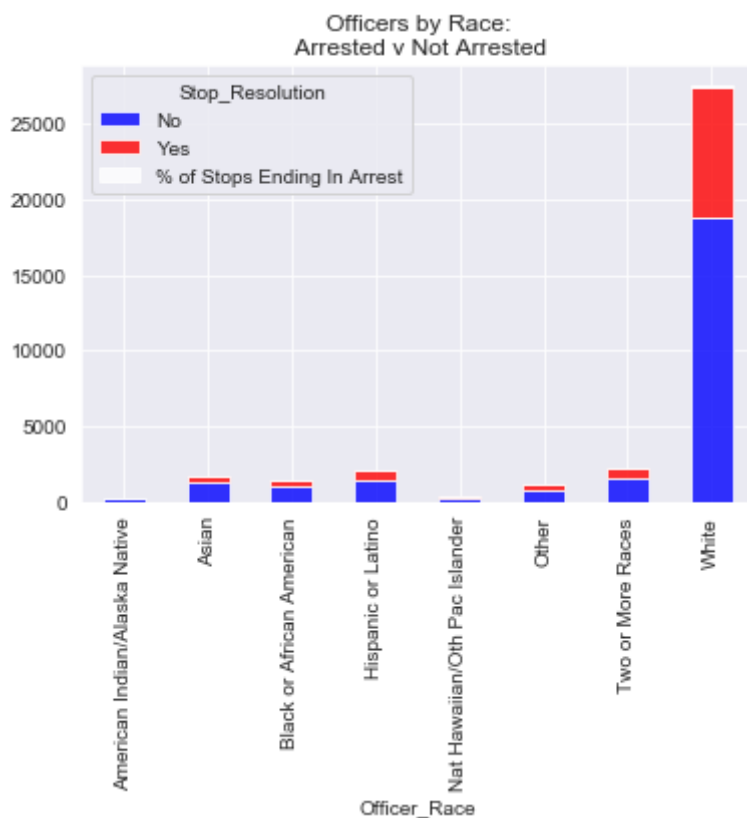
```
In [70]: officer_race = df.groupby(['Officer_Race', 'Stop_Resolution']).Subject_Percentage
officer_race['% of Stops Ending In Arrest'] = (officer_race['Yes'] / (officer_race['No'] + officer_race['Yes'])) * 100

officer_race.reset_index(inplace=True)
officer_race.sort_values('% of Stops Ending In Arrest', ascending=False, inplace=True)

officer_race.plot(kind='bar', stacked=True, title='Officers by Race:\n Arrested v Not Arrested', color=['blue', 'red', 'white'], alpha=0.5)
```

Officer Race

Stop_Resolution	No	Yes	% of Stops Ending In Arrest
Officer_Race			
American Indian/Alaska Native	149	69	31.651376
Asian	1203	446	27.046695
Black or African American	1023	323	23.997028
Hispanic or Latino	1407	653	31.699029
Nat Hawaiian/Oth Pac Islander	246	84	25.454545
Other	767	341	30.776173
Two or More Races	1526	677	30.730822
White	18764	8685	31.640497



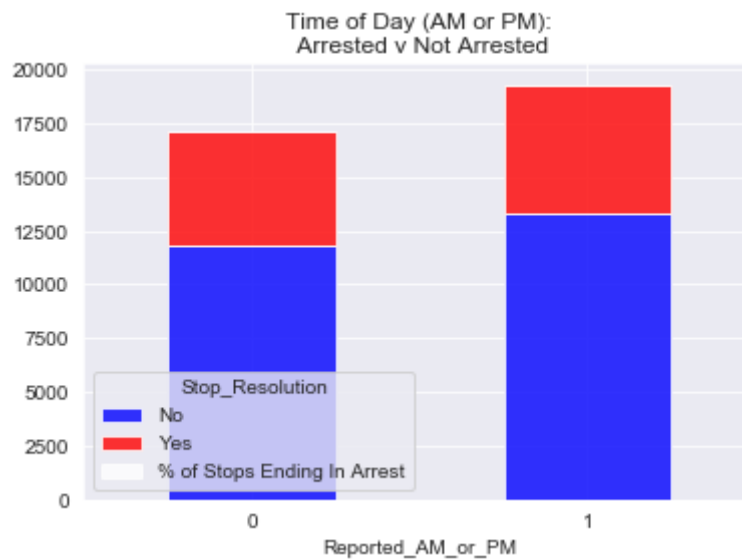
In [71]: *# While white officers make up nearly 3 times the cops of all other race
combined, they were actually in the middle of the pack in terms of
arrests made. Officers of two + races, Hispanic/Latino, and Asian offi
were arresting at around ~25%.*

```
In [72]: am_or_pm = df.groupby(['Reported_AM_or_PM', 'Stop_Resolution']).Reported
am_or_pm['% of Stops Ending In Arrest'] = (am_or_pm['Yes'] / (am_or_pm.s

print('Reported AM or PM\n')
print(am_or_pm)

viz_6 = am_or_pm.plot(kind = 'bar', stacked = True,
                        title = 'Time of Day (AM or PM):\n Ar
                        color = ['blue', 'red','white'], alph
```

0	11821	5286	30.899632
1	13264	5992	31.117574



In [73]: *There are more Terry stops in the PM hours than AM, however the percentages are less than 0.3% apart.*


```
In [74]: hour_of_day = df.groupby(['Reported_Hour', 'Stop_Resolution']).Reported_
hour_of_day['% of Stops Ending In Arrest'] = (hour_of_day['Yes'] / (hour

print('Reported Hour\n')
print(hour_of_day)

viz_7 = hour_of_day.plot(kind = 'bar', stacked = True,
                        title = 'Hour of Day:\n Arrested v No
                        color = ['blue', 'red', 'white'], alph
```

Reported Hour

Stop_Resolution	No	Yes	% of Stops Ending In Arrest
Reported_Hour			
0	1170	517	30.646117
1	1237	587	32.182018
2	1527	606	28.410689
3	1400	565	28.753181
4	766	377	32.983377
5	795	374	31.993157
6	651	320	32.955716
7	679	300	30.643514
8	687	347	33.558994
9	788	332	29.642857
10	1078	409	27.505044
11	1043	552	34.608150
12	776	441	36.236647
13	917	415	31.156156
14	1076	443	29.163924
15	1129	495	30.480296
16	1205	488	28.824572
17	1304	572	30.490405
18	1456	637	30.434783
19	1367	606	30.714648
20	725	448	38.192668
21	925	412	30.815258
22	1100	478	30.291508
23	1284	557	30.255296



```
In [75]: # Looking at the data, arrests ebb and flow depending on time of day,
# with peaks at midnight to 3am and 4pm to 7pm. However the highest
# percentage of arrests made at 11am and 12pm, over 1 percentage point
# away from the next two closest hours at 9am and 8pm.
```

Step 3: Model the Data

```
In [76]: # Now that we've cleaned and visualized the data, now let's fit it into
# model. First starting with converting the target variable to binary,
# OneHotEncoding, then normalizing the data with StandardScaler
```

```
In [77]: df['Stop_Resolution'] = df['Stop_Resolution'].apply(lambda x: 0 if x=='N
df['Stop_Resolution'].value_counts()
```

```
Out[77]: 0    25085
1     11278
Name: Stop_Resolution, dtype: int64
```

```
In [78]: one_hot_df = pd.get_dummies(df)

         = one_hot_df['Stop_Resolution']
one_hot_df.drop('Stop_Resolution', axis=1, inplace=True)

X_train, X_test, y_train, y_test = train_test_split(one_hot_df, y, test_s
```

```
In [79]: scaler = StandardScaler()

scaled_data_train = scaler.fit_transform(X_train)
scaled_data_test = scaler.transform(X_test)

scaled_df_train = pd.DataFrame(scaled_data_train, columns=one_hot_df.col
```

```
In [80]: scaled_df_train.head()
```

```
Out[80]:
```

	Officer_YOB	Reported_Date	Reported_Hour	Reported_AM_or_PM	Subject_Age_Group_1 - 17	Sub
0	0.386363	-0.726015	-1.070069	-1.060617	-0.216547	
1	0.731738	-1.025905	-1.486164	-1.060617	-0.216547	
2	-2.146385	-0.426125	0.178215	0.942847	-0.216547	
3	0.731738	1.373217	1.426499	0.942847	-0.216547	
4	1.307362	0.773437	0.871706	0.942847	-0.216547	

```
In [81]: # Now that the data is OneHotEncoded and normalized, let's fit a model,
# starting with a Logistic Regression model.
```

```
In [82]: logreg = LogisticRegression(fit_intercept=False, C=1e12, solver='libline')
logreg.fit(scaled_data_train, y_train)
```

```
Out[82]: LogisticRegression(C=1000000000000.0, fit_intercept=False, solver='liblinear')
```


```
In [83]: y_hat_train = logreg.predict(scaled_data_train)
y_hat_test = logreg.predict(scaled_data_test)
```

```
In [84]: print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test))
```

Testing Accuracy: 0.5704661075209679

```
In [85]: Looking at the initial model, both the training and the test, both had a
accuracy of 57%. We can definitely make this better using more sophisticated
modeling. Next let's look at a confusion matrix to see where we can
see how the model is performing
```

```
In [86]: plot_confusion_matrix(logreg, scaled_data_test, y_test,
                                cmap=plt.cm.Blues)
plt.grid(False)
plt.show()
```



```
In [87]: # The model is correctly predicting stops ending in no arrest 2791 times
# while incorrectly predicting stops that were arrests as non arrest 897
# times.

# Next let's check the Precision, Recall, Accuracy, and F1 scores.
```

```
In [88]: print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test))
print('Testing Precision: ', precision_score(y_test, y_hat_test))
print('Testing Recall: ', recall_score(y_test, y_hat_test))
print('Testing F1-Score: ', f1_score(y_test, y_hat_test))
```

```
Testing Accuracy:  0.5704661075209679
Testing Precision:  0.3788005578800558
Testing Recall:    0.6022172949002217
Testing F1-Score:  0.4650684931506849
```


```
In [89]: # Not too bad for a first start except for the precision.
# This time, we'll fit a K Nearest-Neighbors model to see if the data
# works better in that context.
```

```
In [90]: clf = KNeighborsClassifier()
clf.fit(scaled_data_train, y_train)
y_hat_train = clf.predict(scaled_data_train)
y_hat_test = clf.predict(scaled_data_test)
```

```
In [91]: print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test))
print('Testing Precision: ', precision_score(y_test, y_hat_test))
print('Testing Recall: ', recall_score(y_test, y_hat_test))
print('Testing F1-Score: ', f1_score(y_test, y_hat_test))
```

```
Testing Accuracy:  0.6467757459095284
Testing Precision:  0.3969816272965879
Testing Recall:    0.2682926829268293
Testing F1-Score:  0.3201905265943371
```

```
In [92]: plot_confusion_matrix(clf, scaled_data_test, y_test,
                               cmap=plt.cm.Blues)
plt.grid(False)
plt.show()
```



In [93]: *The accuracy and precision went down slightly in comparison to the Logistic Regression output. A confusion matrix shows that it has only predicted '0', not a great sign. Since KNN has a default "neighbors" of 5, let's use GridSearchCV to see if we can find a the best number of neighbors and run that model one more time.*

Source of following code: Eijaz Allibhai (<https://towardsdatascience.com>)

```
In [106]: clf2 = KNeighborsClassifier()
param_grid = {'n_neighbors': np.arange(1,25)}
clf_gscv = GridSearchCV(clf2, param_grid, cv=5)
clf_gscv.fit(scaled_data_train, y_train)
```

```
Out[106]: GridSearchCV(cv=5, estimator=KNeighborsClassifier(),
                      param_grid={'n_neighbors': array([ 1,  2,  3,  4,  5,  6,
 7,  8,  9, 10, 11, 12, 13, 14, 15, 16, 17,
 18, 19, 20, 21, 22, 23, 24])})
```

```
In [95]: clf_gscv.best_params_
```

```
Out[95]: {'n_neighbors': 24}
```

In [96]: *# Let's rerun the KNN model to see exactly how this 24 is the best neigh*

```
In [97]: clf = KNeighborsClassifier(n_neighbors=24)
clf.fit(scaled_data_train, y_train)
y_hat_test = clf.predict(scaled_data_test)
```

```
In [98]: print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test))
print('Testing Precision: ', precision_score(y_test, y_hat_test))
print('Testing Recall: ', recall_score(y_test, y_hat_test))
print('Testing F1-Score: ', f1_score(y_test, y_hat_test))
```

```
Testing Accuracy:  0.6837618589302901
Testing Precision:  0.4403183023872679
Testing Recall:    0.07361419068736141
Testing F1-Score:  0.12613981762917933
```

In [99]: *# Accuracy and precision are up but recall and F1 are not great.*
Let's move on to XGBoost

```
In [100]: xgbc = XGBClassifier()
xgbc.fit(scaled_data_train, y_train)
y_hat_test = xgbc.predict(scaled_data_test)
```



```
In [101]: print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test))
print('Testing Precision: ', precision_score(y_test, y_hat_test))
print('Testing Recall: ', recall_score(y_test, y_hat_test))
print('Testing F1-Score: ', f1_score(y_test, y_hat_test))
```

```
Testing Accuracy:  0.6902241165956277
Testing Precision:  0.5098039215686274
Testing Recall:    0.023059866962305987
Testing F1-Score:  0.044123886296139156
```

```
In [102]: # Accuracy and Precision are up, but everything else is worse.
# Let's try tuning XGBoost using GridSearchCV
# to see if we get better results.
```

```
In [109]: param_grid = {
    'learning_rate': [0.1, 0.2, 0.3, 0.4, 0.5],
    'max_depth': [6],
    'min_child_weight': [1, 2, 3, 4, 5],
    'subsample': [0.5, 0.7, 0.9],
    'n_estimators': [100],
}
```

```
In [110]: xgbc = XGBClassifier()
grid_xgbc = GridSearchCV(xgbc, param_grid, scoring='accuracy', cv=None,
grid_xgbc.fit(scaled_data_train, y_train)

best_parameters = grid_xgbc.best_params_
```

```
In [111]: print(best_parameters)

{'learning_rate': 0.1, 'max_depth': 6, 'min_child_weight': 2, 'n_estimators': 100, 'subsample': 0.9}
```

```
In [112]: xgbc = XGBClassifier(n_estimators=100, max_depth=6, learning_rate=0.1,
    min_child_weight=2, subsample=0.9)
xgbc.fit(scaled_data_train, y_train)
y_hat_test = xgbc.predict(scaled_data_test)
```

```
In [113]: print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test))
print('Testing Precision: ', precision_score(y_test, y_hat_test))
print('Testing Recall: ', recall_score(y_test, y_hat_test))
print('Testing F1-Score: ', f1_score(y_test, y_hat_test))
```

```
Testing Accuracy:  0.690774095971401
Testing Precision:  0.5081081081081081
Testing Recall:    0.08337028824833703
Testing F1-Score:  0.14323809523809525
```

```
In [114]: # This is the highest our accuracy has been but barely. Was the 30
# minutes of loading this GridSearch worth it?
# Let's move to a decision tree and see if we can get better than
# 69% accuracy.
```

```
In [115]: SEED=1
dtc = DecisionTreeClassifier(criterion='entropy', random_state=SEED)
dtc.fit(scaled_data_train, y_train)
```

```
Out[115]: DecisionTreeClassifier(criterion='entropy', random_state=1)
```

```
In [116]: y_hat_test = dtc.predict(scaled_data_test)

print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test))
print('Testing Precision: ', precision_score(y_test, y_hat_test))
print('Testing Recall: ', recall_score(y_test, y_hat_test))
print('Testing F1-Score: ', f1_score(y_test, y_hat_test))
```

```
Testing Accuracy:  0.6187267977450845
Testing Precision:  0.38748913987836664
Testing Recall:    0.39556541019955654
Testing F1-Score:  0.39148562650866797
```

```
In [117]: SEED=1
dtc = DecisionTreeClassifier(criterion='gini', random_state=SEED)
dtc.fit(scaled_data_train, y_train)
```

```
Out[117]: DecisionTreeClassifier(random_state=1)
```

```
In [118]: y_hat_test = dtc.predict(scaled_data_test)

print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test))
print('Testing Precision: ', precision_score(y_test, y_hat_test))
print('Testing Recall: ', recall_score(y_test, y_hat_test))
print('Testing F1-Score: ', f1_score(y_test, y_hat_test))
```

```
Testing Accuracy:  0.6147394472707274
Testing Precision:  0.38580375782881005
Testing Recall:    0.4097560975609756
Testing F1-Score:  0.3974193548387097
```

```
In [119]: # The difference between 'gini' and 'entropy' were basically non existen
# Let's try another GridSearch to see instead of manually changing paran
# Hopefully we can find the best decision tree.
```

```
In [121]: dtc_param_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 2, 3, 4, 5],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 3, 4, 5]
}

dtc_grid = GridSearchCV(dtc, dtc_param_grid, cv=3, return_train_score=True)

dtc_grid.fit(scaled_data_train, y_train)
```

```
Out[121]: GridSearchCV(cv=3, estimator=DecisionTreeClassifier(random_state=1),
    param_grid={'criterion': ['gini', 'entropy'],
    'max_depth': [None, 2, 3, 4, 5],
    'min_samples_leaf': [1, 2, 3, 4, 5],
    'min_samples_split': [2, 5, 10]},
    return_train_score=True)
```

```
In [122]: dtc_grid.best_params_
```

```
Out[122]: {'criterion': 'gini',
    'max_depth': 2,
    'min_samples_leaf': 1,
    'min_samples_split': 2}
```

```
In [123]: dtc = DecisionTreeClassifier(criterion='gini', max_depth=2,
    min_samples_leaf=1, min_samples_split=2,
    random_state=SEED)

dtc.fit(scaled_data_train, y_train)
```

```
Out[123]: DecisionTreeClassifier(max_depth=2, random_state=1)
```

```
In [124]: y_hat_test = dtc.predict(scaled_data_test)

print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test))
print('Testing Precision: ', precision_score(y_test, y_hat_test))
print('Testing Recall: ', recall_score(y_test, y_hat_test))
print('Testing F1-Score: ', f1_score(y_test, y_hat_test))
```

```
Testing Accuracy:  0.6899491269077409
Testing Precision:  0.0
Testing Recall:    0.0
Testing F1-Score:  0.0
```

```
/Users/Garseed/anaconda3/envs/learn-env/lib/python3.6/site-packages/sklearn/metrics/_classification.py:1248: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

```
In [125]: # It seems as though we were unable to get enough data for testing
# metrics outside of Accuracy, probably due to the low samples.
```

```
In [126]: # After running 7 models, the highest accuracy achieved was from an  
# XGBoost model with the best parameters found through GridSearchCV,  
# and even at that it was only able to be 69% accurate.
```